

HPC Project Report

Leonardo Antonio Riola, Davide Donati, Giacomo Girotto

1 Assignment 1

1.1 Introduction and objectives

The goal of this assignment is to implement a simulator for the propagation of a damped wave, governed by the following partial differential equation:

$$\frac{\partial^2 u}{\partial t^2} + \gamma \frac{\partial u}{\partial t} = c^2 \nabla^2 u \quad (1)$$

where $u(x, y, t)$ is the wave height on a two-dimensional plane over time, γ is the damping coefficient, and c is the propagation speed. The simulator computes u at every point (x, y) as t advances, with the final goal of producing an MP4 video of the simulation.

1.2 Implementation

The continuous model is discretised on a regular grid: space is divided into points spaced by Δx and Δy , and time into steps of Δt . Each grid point is identified by indices (i, j) and each instant by an index t , so that $u(x, y, t)$ becomes a matrix of values $u_{i,j}^t$. The derivatives in the equation are approximated with **central differences method** [14], so that the next wave state $u_{i,j}^{t+1}$ is:

$$u_{i,j}^{t+1} = \frac{1}{\gamma \Delta t + 2} \left[c^2 \Delta t^2 \left(\frac{u_{i+1,j}^t - 2u_{i,j}^t + u_{i-1,j}^t}{\Delta x^2} + \frac{u_{i,j+1}^t - 2u_{i,j}^t + u_{i,j-1}^t}{\Delta y^2} \right) + \gamma u_{i,j}^{t-1} \Delta t - 2u_{i,j}^{t-1} + 4u_{i,j}^t \right] \quad (2)$$

The temporal evolution is handled with three $M \times M$ matrices, **u_past**, **u** and **u_fut**, dynamically allocated with **malloc** and representing the wave state at times $t-1$, t and $t+1$ respectively. All elements are initialised to zero except for a single impulse at the centre of the grid $(M/2, M/2)$, used to start the perturbation. The impulse amplitude, propagation speed and damping coefficient used throughout this work are:

Impulse	Propagation speed c	Damping coefficient γ
-49	0.33	0.187

The core of this assignment is the parallelization of the update loop with **OpenMP**. A **#pragma omp parallel** region is opened once before the time loop, and the two nested spatial loops are distributed across threads with **#pragma omp for schedule(static) collapse(2): schedule(static)** splits the iterations into equal-sized chunks assigned before execution [9], while **collapse(2)** fuses the i and j loops into a single $M \times M$ iteration space [6], giving the scheduler far more, finer-grained work units to balance across cores. Each thread thus computes $u_{i,j}^{t+1}$ for a disjoint set of coordinates, independently and concurrently with the others.

Once **u_fut** has been fully computed for the current timestep, it must be written to a PGM file; this is done inside a **#pragma omp single** block, so that only one thread performs the write while the rest wait at the implicit barrier. To avoid the cost of deep-copying entire matrices between timesteps, the temporal advance is implemented as a three-way pointer rotation (**u_past** takes the address of **u**, **u** takes that of **u_fut**, and **u_fut** is recycled for the next iteration), performed by the same thread that just wrote the file. This keeps the pointer swap race-free and avoids any extra synchronization beyond the barrier already required for the write.

1.2.1 Identified inefficiencies and corrective transformations

All profiling and testing in this section was performed using the `Edu_skylake` partition (4 nodes, 32 physical and logical cores) [1] with 32 threads and a workload size of $M = 1000$. The aim was to identify the main bottlenecks and improve the source code.

We profiled the initial implementation with **Intel VTune Profiler**. *Hotspots* analysis identifies which functions in an application consume the largest share of execution time, using low-overhead sampling to build a ranked list of hotspot functions as a starting point for optimization [3]. *Threading* analysis instead evaluates how efficiently the application uses the available processor cores, exposing inefficiencies in threading-runtime usage and contention at synchronization points that leave threads idle and prevent full processor utilization, summarized through its *Effective CPU Utilization* metric [4]. Combining the two analyses lets us separate genuine computation time from time lost to idle spinning at synchronization points.

Hotspot analysis showed that `gomp_team_barrier_wait_end` and `fwrite` accounted for 79.4% and 9.7% of total CPU time respectively, against only 2.2% for the matrix computation itself: the bottleneck was therefore not the stencil, but excessive synchronization and I/O overhead. Three successive transformations addressed this, each verified either via VTune or direct code inspection.

Improvement 1: buffered frame output. Frames were originally written pixel by pixel with `fwrite()` inside a single-threaded `#pragma omp single` block. We replaced this with an in-memory buffer, filled in parallel via a dedicated `#pragma omp for` and flushed to disk with one `fwrite()` call, cutting I/O calls from one per pixel to one per frame. This did introduce a new barrier, however, since the buffer-fill loop is a separate worksharing construct from the stencil loop.

Improvement 2: loop fusion. Since the stencil and the buffer-fill iterate over the exact same (i, j) range, we merged them into a single `#pragma omp for`, writing each scaled pixel immediately after computing `u_fut[i][j]`. This removes the extra barrier introduced above, with no loss of parallel work.

Improvement 3: parallel initialization and scheduling policy. The matrix initialization, originally performed sequentially by a single thread, was parallelized with `#pragma omp parallel for`. By Amdahl’s law, this serial fraction would otherwise cap the achievable speedup regardless of how well the main loop scales. On a multi-socket compute node, each CPU socket has its own directly-attached memory bank: cores access memory local to their own socket faster than memory attached to another socket, a design known as NUMA (Non-Uniform Memory Access) [8]. Under the default *first-touch* policy, a physical page is assigned to the socket of whichever thread writes to it first, we therefore gave the initialization loop the same schedule as the stencil loop, so that each thread’s first-touch region matches the region it repeatedly accesses afterwards. Profiling the complete pipeline (including I/O) showed, however, that `schedule(guided)` outperforms `schedule(static)` in practice ¹: under `guided` scheduling, chunks are handed out to threads on request rather than fixed in advance, with each chunk’s size proportional to the number of remaining iterations and shrinking toward the specified minimum (1 as default) as the loop progresses [6]. This dynamic redistribution better compensates for the load imbalance introduced by the single-threaded frame-writing step, so `guided` was retained for the final implementation. As shown later in the testing section, isolating the computation from I/O reverses this result, confirming that NUMA locality does matter once the synchronization overhead of I/O is removed.

After these three improvements, VTune Hotspot analysis showed the CPU time spent in `gomp_team_barrier_wait_end` dropping to 57.5%, while time spent in the matrix computation itself rose to 15.7%; I/O (`fwrite`) time fell below 0.3% ². The Threading analysis, however, revealed that the program was still using only 23% of the available CPUs, for an average effective utilization of 7.362 out of 32 threads: a clear sign that, at $M = 1000$, the workload was simply too small for the

¹See `output_1920835.txt` (static) and `output_1920842.txt` (guided)

²See `output_1920920.txt`

number of threads requested. This motivated the next round of tests, which vary the workload size and other parameters to see how resource utilization changes.

1.3 Testing and Results

The remaining tests were run on the `Edu_sapphire` partition (2 nodes, 48 physical and 96 logical cores each); each node is dual-socket, with 24 physical cores per socket [1].

1.3.1 Sweep of workload size




Figure 1: Execution time and true CPU efficiency scaling with varying grid sizes (M), using 96 threads.

In order to investigate the architectural limits of the system, a workload sweep test was conducted using the maximum logical capacity of the compute node (96 threads), varying the grid size M from 1,000 to 20,000. Figure 1 illustrates the execution time (extracted without the profiling overhead) vs the True Effective CPU Utilization, extracted from the VTune threading analysis. The utilization curve reveals three distinct operational phases. For small grids ($M = 1000$ to 2500), the system suffers from severe *thread starvation*, where the computational workload is too small to justify the distribution across 96 threads and two NUMA sockets, resulting in an initial efficiency plateau around 21%, dominated by synchronization overhead, barrier wait times and I/O operations. As M increases ($M = 5000$ to 15000), the computational density grows quadratically ($O(M^2)$) the useful computational work begins to outweigh the fixed parallelization overhead, driving the CPU utilization upward.

However, the efficiency strictly plateaus at approximately 36% for massive grids ($M \geq 15000$). This hard limit is the superposition of three architectural constraints: the strict sequential bottleneck of the disk I/O operations (Amdahl's Law), the saturation of the memory bandwidth (memory wall), and the effects of Hyper-Threading. Regarding the latter, since the node possesses only 48 physical cores, through Hyper-Threading each core exposes two logical processors to the OS, which share the same execution units, caches and memory bandwidth [5]. So, attempting to utilize 96 threads for a memory-bound stencil operation forces logical threads to heavily compete for the same physical cache resources, preventing the utilization from approaching 100%. Regarding the absolute execution times, the data perfectly validates the theoretical $O(M^2)$ spatial complexity of the 2D grid simulation. For instance, doubling the grid dimension from $M = 10000$ to $M = 20000$ effectively quadruples the computational domain and the execution time scales from 33.0 seconds to 123.6 seconds, achieving a 3.74x increase that is remarkably close to the theoretical 4x factor. Conversely, in the smaller grid range ($M \leq 2500$), the constant overhead associated with managing 96 concurrent threads, along with the fixed costs of implicit OpenMP barriers, heavily dominates the actual execution time, confirming the thread starvation phenomenon observed in the efficiency curve.

1.3.2 Sweep of number of threads at $M = 15000$




Figure 2: Speedup and parallel efficiency with varying number of threads.

We then tested the program with a fixed workload size of $M = 15000$ varying the number of threads and collected the **Speedup** and **Parallel Efficiency** (Speedup/Number of threads) parameters. Figure 2 shows two distinct regimes. For $P \leq 32$, all threads run on independent physical cores; the efficiency loss is explained by the OpenMP barriers cost, which grows with team size. For $t = 64$ and $t = 96$, thread count exceeds the number of physical cores (48), so Hyper-Threading is used from that point onward.

We repeated the same test with the I/O operations disabled, so that the program only computes the matrix updates.³




Figure 3: Speedup and parallel efficiency with varying number of threads, I/O disabled.

As Figure 3 shows, both metrics improve significantly: at $t = 96$ speedup more than doubles, from $11.1\times$ to $25.1\times$, and parallel efficiency grows from 11.6% to 26.1% . This confirms that a substantial share of the overhead seen in the full version came from the I/O operations and their associated synchronization barriers, rather than from the stencil computation itself.

Finally, we replaced `schedule(guided)` with `schedule(static)` in both the initialization and the stencil loop, to restore the NUMA alignment mentioned earlier.⁴ Once thread count exceeds 24, work is necessarily split across both sockets, and accesses to memory owned by the other socket incur the added latency of the inter-socket interconnect.

³See output_1925400.txt, output_1925398.txt, output_1925396.txt, output_1925394.txt, output_1925399.txt, output_1925397.txt, output_1925395.txt, output_1925393.txt in folder "results_noio"

⁴See output_1928239.txt, output_1928222.txt, output_1928220.txt, output_1927019.txt, output_1927027.txt, output_1927036.txt in folder "results_noio-static"




Figure 4: Speedup and parallel efficiency with varying number of threads, I/O disabled, `schedule(static)`.

As shown in Figure 4, the static schedule substantially mitigates this effect: the efficiency drop between $t = 16$ and $t = 32$ (the exact point where the 24-core NUMA boundary is crossed) is roughly halved (from -19.3 to -10.0 percentage points relative to **guided**), and efficiency remains markedly higher at $t = 32$ (43.7% vs. 29.9%), $t = 64$ (45.4% vs. 27.0%, a 68% relative gain) and $t = 96$ (38.3% vs. 26.1%), with speedup at $t = 96$ rising from $25.1\times$ to $36.8\times$. This is consistent with the static schedule preserving the same thread-to-row mapping established during initialization, which keeps most memory accesses local to each thread’s own socket even as the workload spans both NUMA domains.

2 Assignment 2

2.1 Introduction

Following the initial wave simulator developed with OpenMP, this second assignment extends the application to execute multiple concurrent simulations on an HPC system using a hybrid MPI and OpenMP programming model. The goal is to study multi-wave interaction and interference dynamics under different physical conditions. While OpenMP continues to handle data parallelism by accelerating the 2D stencil computations across local CPU cores, MPI is introduced to manage task parallelism. Specifically, the system executes three distinct simulation scenarios simultaneously within a single execution run:

- **Simulation 1 (Rank 0):** Reproduces the baseline single-wave scenario initiated by a single central impulse at coordinates $(M/2, M/2)$ at $t = 0$.
- **Simulation 2 (Rank 1):** Models wave superposition and interference phenomena by initiating two simultaneous impulses at different spatial locations at $t = 0$.
- **Simulation 3 (Rank 2):** Introduces time-delayed interference dynamics by generating a second impulse at a spatial location after a predetermined time-shift ($t = t_{start}$).

By decoupling these independent workloads across separate MPI processes, the hybrid architecture maximizes computational throughput, since each rank advances its own simulation through the full iterative loop without ever needing to exchange data with the other two ranks during execution.

2.2 Implementation

Building upon the code developed in Assignment 1, the implementation extends the simulator to handle multiple concurrent runs. The core numerical scheme used to advance the wave field at each iteration, together with the mapping function converting field values into grayscale images (`write_matrix`), remain identical to the first version, ensuring consistency in wave physics and visualization.

The execution starts by initializing the MPI environment with **MPI_Init**, retrieving the total number of processes with **MPI_Comm_size**, and identifying the rank ID of each process with **MPI_Comm_rank**. To run the three required simulations simultaneously, the program verifies that at least 3 processes are allocated. Task parallelism is managed through conditional branching based on the process rank, where each rank specifies its scenario parameters before calling a dedicated **run_simulation** function that encapsulates the core iterative loop advancing the simulation in time.

While data parallelism within each process is handled by OpenMP, the execution incorporates rank-specific logic. Thread synchronization constructs, such as the **pragma omp single** directive, are carefully utilized within the main loop to safely manage tasks that must be executed by a single thread. These tasks include array pointer swapping, injecting the delayed impulse, and saving frames, and confining them to a single thread prevents race conditions without halting the entire team of threads unnecessarily.

2.3 Test and optimization

Two main optimizations were applied to the algorithm on C. The first is about memory management: each MPI process independently allocates its local grid arrays (**u_past**, **u**, and **u_fut**) through a custom **matrix_alloc** function that guarantees contiguous memory. Even though the three ranks are fully decoupled and never communicate, contiguous allocation remains essential for the local OpenMP threads, as it maximizes spatial locality and hardware prefetching efficiency during the 2D stencil computations, mitigating the RAM contention generated by three concurrent processes running side by side.

The second optimization tries to solve the problem of thread synchronization. Operations that must be executed by a single thread, like pointer swapping, frame saving, and the injection of the time-delayed impulse for Simulation 3, require a **pragma omp single** construct. To minimize its cost, the condition **if (rank == 2 && t == n_start)** was placed *outside* the **pragma omp single** block rather than inside it. As a result, the synchronization barrier is only reached by the threads of Rank 2 at the exact time step in which the impulse must be injected; during every other iteration, and for the other two ranks at all times, the threads bypass the barrier entirely, avoiding unnecessary stalls of the whole thread team.

With these optimizations in place, the three ranks were profiled with 8 OpenMP threads each, using the `edu.sapphire` partition as in Assignment 1. Execution times in 5, showed no significant variation between ranks, confirming a balanced workload despite their different initial conditions and branching logic. However, when compared against the Assignment 1 baseline, a severe performance degradation appeared as the matrix size M increased.




Figure 5: Execution time with 8 threads per rank by swapping the M size.

This slowdown could not be attributed to the computational cost of the hybrid setup itself. Instead, it was caused by an Input/Output bottleneck: all three MPI processes attempt to format and write **.pgm** frames to the shared filesystem simultaneously, generating heavy hardware contention that

stalls the CPU threads and inflates the overall execution time [7]. To isolate the purely computational performance of the architecture, we have disabled the operation of writing to disk, and the profiling campaign was repeated while varying M with 8 OpenMP threads per MPI task.

In a first configuration, all three ranks, 24 threads in total, were confined to a single CPU socket of `edu_sapphire`, without strict thread pinning, in order to evaluate the shared-memory behavior among the processes. For each rank, both a hotspots and a threading analysis were collected independently using Intel VTune Profiler.




Figure 6: Execution time and effective CPU utilization with 24 threads in one socket, swapping the M .

As shown in Figure 6, removing the I/O bottleneck alone produces a drastic reduction in execution time, from over 700 seconds down to approximately 61 seconds for $M = 10000$. However, Intel VTune profiling revealed a significant hardware-level anomaly: the effective CPU utilization is highly asymmetric across ranks, with Rank 1 showing a markedly lower efficiency than Rank 0 and Rank 2, even though all three processes perform the exact same amount of computation. Since the three ranks were running within the same socket without strict core pinning, the Linux scheduler occasionally migrated threads from busier cores to idle ones in an attempt to balance the load. Every time a thread resumes execution on a new core, it finds none of the grid data it needs in the local cache, forcing the processor to fetch it from main RAM, an operation orders of magnitude slower than a cache hit, which effectively stalls the core while it waits for the data [2, 7]. This phenomenon disproportionately affected Rank 1’s threads, and the resulting memory stalls are the direct explanation for its degraded effective CPU utilization relative to the other two processes.

To validate this hypothesis about memory on a topology where socket boundaries could be controlled more explicitly, a second experiment was conducted on the `edu_skylake` partition, whose nodes expose two physical sockets of 16 cores each. Three tasks of 8 threads each were requested via `-ntasks=3`, without explicit process binding, relying on OpenMPI’s default placement policy. Since each socket can host at most 16 threads, the expectation was that the three ranks would naturally split across the two sockets in a 2-and-1 arrangement rather than being packed together. The profiling results confirm this: Rank 0 and Rank 2 ended up sharing the same socket, while Rank 1 was cyclically placed alone on the second socket.

The results in Figure 7 confirm the memory hypothesis: with Rank 0 and Rank 2 sharing 16 physical cores on one socket and Rank 1 alone on the other socket, Rank 1 completed its stencil computations markedly faster than the other two, like 88 seconds versus 150 seconds at $M = 15000$.

This same test also reproduces, on a different hardware topology, the effective CPU utilization drop observed for the faster rank: because `MPI_Finalize` is a global synchronization point, Rank 1 finishes its computation early and then idles in a spin-wait state until Ranks 0 and 2 complete. Since VTune averages effective CPU utilization over the entire process lifetime, this prolonged idle period drags down Rank 1’s recorded efficiency, even though it is the fastest of the three ranks to finish.

To further isolate the effect of memory locality identified in the previous experiment, a final test was conducted by running only Simulation 2 and Simulation 3 concurrently, each assigned to an




Figure 7: Execution time and effective CPU utilization with Rank 1 isolated on a separate socket of edu_skylake partition, swapping M.

entirely separate physical node on edu_sapphire; Simulation 1 was excluded, since it reproduces the same single-wave experiment already characterized in Assignment 1. As this test was launched as an independent two-process MPI job, the ranks were renumbered accordingly, so Rank 0 and Rank 1 in the figure below correspond to Simulation 2 and Simulation 3 respectively. Unlike the previous setup, where isolation was achieved at the socket level within a single shared node, this configuration grants each rank exclusive access not only to its own socket, but to the entire node’s memory subsystem, with no possibility of interference from the other rank. The number of OpenMP threads per rank was varied from 1 to 48 to additionally assess parallel scalability under this fully isolated condition.




Figure 8: Execution time and effective CPU utilization with Rank 0 and Rank 1 on separate physical nodes, sweep the number of threads. The two diamond markers at 8 threads report the execution times measured in the previous experiments under partial isolation, for direct comparison.

As shown in Figure 8, this configuration produces a further reduction in execution time at 8 threads compared to the previous experiments: execution time drops from approximately 150 seconds per process, measured when three ranks shared the same socket, and approximately 88 seconds, measured when one rank was isolated on a separate socket within the same node, down to roughly 57 seconds (58.9s for Rank 0, 55.9s for Rank 1) when each rank is granted an entire dedicated node.

A second, equally important observation is that Rank 0 and Rank 1 now exhibit virtually identical execution times and effective CPU utilization at every thread count, in sharp contrast to the asymmetry observed in the single-node, two-socket experiment. Since neither rank finishes significantly ahead of the other, neither is forced into a prolonged idle spin-wait on **MPI_Finalize**, and the CPU utilization metric is no longer artificially diluted for the faster process, indirectly confirming that the asymmetry seen earlier was indeed caused by unequal waiting time rather than unequal computational

work.

Finally, effective CPU utilization increases steadily with the thread count, from about 1% at a single thread to approximately 48% at 48 threads, reflecting the growing fraction of the node’s 96 logical cores actively engaged in computation. The best configuration observed across this campaign is therefore $M = 15000$ with 48 OpenMP threads per rank and one dedicated physical node per rank, which minimizes execution time to approximately 29 seconds, while making full use of the node’s 48 physical cores without contending for shared memory bandwidth.

3 Assignment 3

3.1 Introduction

This assignment extends the damped wave simulator from parts 1–2 with color. The physics, discretisation, and MPI/OpenMP structure are unchanged: three MPI ranks each run one independent scenario (a single impulse; two impulses starting together; two impulses with the second delayed to $N/7$), each parallelised with OpenMP. What changes is the output stage: each timestep’s matrix is sent to the GPU, mapped to an RGB triplet per pixel by a CUDA kernel, and saved as a P3 (ASCII) PPM frame instead of a grayscale PGM. Bright, saturated colors mark the wave’s extrema, colors fade along the rising/falling fronts, and calm (near-zero) regions are rendered white.

3.2 Architecture

The three MPI ranks [10] map naturally onto the three required scenarios: each rank owns a complete, independent grid and runs one scenario end to end, so the only collective call needed is a final `MPI_Allreduce` on success/failure. Within each rank, `compute_next` (the same 5-point Laplacian stencil as parts 1–2 – each point’s next value is computed from itself and its 4 orthogonal neighbors via central differences) is parallelised with OpenMP [9], using `schedule(static)` for the same NUMA first-touch reason argued in part 1: with a fixed thread-to-row mapping, the memory each thread first writes (and therefore where the OS physically places it) matches the memory that same thread repeatedly accesses on every subsequent call. After each timestep, the rank sends its matrix to the GPU [12]: a kernel normalizes each value against the frame’s own peak magnitude (found via a GPU-side reduction [13]), maps it to a color (white near zero, fading to bright red/orange for positive values and blue for negative, saturating at the wave’s extrema), and the result is written to disk as an ASCII PPM frame while the original scalar matrix continues to drive the simulation. Figure 9 shows one rendered frame. Because the job requests a single GPU (`--gres=gpu:1`) for all three ranks, the GPU is a shared resource – a design point this report returns to in Section 3.8.




Figure 9: A rendered frame (sim1, late in the simulation): bright extrema, fading fronts, white calm regions.

3.3 Cluster hardware

All experiments ran on Politecnico di Torino’s HPC@PoliTO cluster (Legion), specifically its EDU partitions – reserved for coursework, as opposed to the open-access or research-group-restricted partitions also present on the cluster. Of the four EDU partitions, only two have GPUs: `edu_a40` and `edu_v100`, confirmed live via `sinfo -o "%P %D %G"` (the other two, `edu_skylake` and `edu_sapphire`, are CPU-only). `edu_a40` nodes have 2 physical CPU sockets, 24 physical cores per socket (48 physical cores/node), doubled to 96 logical cores/node by SMT (2 hardware threads per core) – confirmed via `sinfo`’s socket/core/thread fields, since the cluster’s own documentation lists only the 48 physical cores – and 4×A40 GPUs; `edu_v100` nodes have 32 cores and 4×V100 GPUs. Both are used in this report: `edu_a40` for all main results, and `edu_v100` as a cross-hardware check in Section 3.6.

3.4 Metrics and the grid-size sweep

Beyond wall-clock time, the program times two phases directly: `color_time` (GPU colorization + PPM write, per frame) and `compute_time` (the OpenMP stencil update). `color_time` is further split into five sub-phases – H2D copy, kernel, D2H copy, ASCII formatting, and disk write – introduced in Section 3.5. Figure 10 sweeps grid size M from 128 to 4096 (`edu_a40`, 4 threads, 50 steps). Execution time approaches ideal $O(M^2)$ scaling as M grows: $128 \rightarrow 256$ costs only $2.1\times$ (fixed per-run overhead still dominant at small M), while $2048 \rightarrow 4096$ costs $4.1\times$ – essentially the ideal quadratic ratio. The right panel shows GPU colorize+write dominating total time throughout: at $M = 4096$ it is $\sim 97.6\%$ of total time, versus `compute_next` at $\sim 2.4\%$. Both phases are $O(M^2)$, so this ratio stays roughly constant across the whole sweep ($\sim 97.6\text{--}97.7\%$ from $M = 256$ upward; only the smallest grid, $M = 128$, deviates, where fixed per-run overhead is large enough to matter).

We capped the sweep at $M = 4096$ rather than matching part 1’s $M = 15000$: at $M = 4096$ an ASCII PPM frame is already $\sim 100\text{--}130\text{MB}$, so scaling to $M = 15000$ – $(15000/4096)^2 \approx 13\times$ the pixels – would push single frames into the multi-gigabyte range across 50 steps and 3 ranks, for a sweep whose purpose (showing execution time approach ideal $O(M^2)$ scaling and locating the point where GPU colorize+write comes to dominate) was already unambiguous well before $M = 4096$.




Figure 10: Execution time and its GPU/OpenMP split, varying grid size M (4 threads, 3 ranks sharing 1 GPU).

3.5 Phase breakdown: is the ASCII formatting worth parallelizing?

Figure 11 breaks `color_time` into its five sub-phases. At $M = 4096$ (averaged over 3 ranks): `write` (the `fwrite` call) is $\sim 94.6\%$ of `color_time`, `format` (the ASCII byte-to-digit conversion) is $\sim 4.3\%$, and H2D copy + kernel + D2H copy combined are $\sim 1\%$. The same lopsided pattern holds at every grid size swept, not only at $M = 4096$.

This directly answers a question we considered early on: is it worth parallelizing the ASCII-formatting step with OpenMP? It is not – formatting is only $\sim 4\%$ of `color_time`, so even a theoretical

zero-overhead parallelization would save a few percent of total time at best, not worth the added complexity (the output is variable-width per pixel, so parallel writers would need either a prefix-sum pass to compute offsets first, or a switch to fixed-width padding). We measured before optimizing, and the measurement said no.




Figure 11: colorize+write split into its 5 phases vs. grid size. Left: full stack (write’s dominance is the point, not log-scaled). Right: same data with write excluded, zoomed in.

3.6 Where does the write time actually go? Nsight Systems, cross-checked on two GPUs

Nsight Compute (kernel-level occupancy/throughput metrics) was attempted but is unavailable on this cluster: it returns `ERR_NVGPUCTRPERM` – GPU performance-counter access is admin-only, a driver-level restriction no user-side flag can work around. Nsight Systems was used instead, and it revealed something our own coarse `write_s` timer could not: that timer wraps `fopen+fwrite+fclose` as one block, and Nsight Systems’ OS Runtime Summary lets us split it apart. On `edu_a40`: `fwrite` itself is cheap ($\sim 0.6\%$ of total OS-runtime time), while `fclose` is $\sim 18.3\%$ – roughly $30\times$ more expensive than the write call it follows. The two largest categories by far, `epoll_wait` ($\sim 52\%$) and `poll` ($\sim 25\%$), we initially guessed reflected ranks idling on MPI’s progress engine. A follow-up run with MPI event tracing enabled rules that out directly: summed across `MPI_Init`, `MPI_Finalize`, and the one real collective (`MPI_Allreduce`), application-level MPI time is ~ 0.56 s per rank – under 2% of the combined `epoll_wait+poll` time (~ 28.4 s). `MPI_Allreduce` itself is 13.7 ms. The better-supported explanation, from a companion VTune Threading run on the same job: its top wait objects include InfiniBand verbs streams (`/dev/infiniband/uverbs0`), suggesting OpenMPI’s UCX transport stack is initialized and idly polling in the background for the whole process lifetime regardless of how little it is actually used – not the application’s own collective calls, but the runtime’s transport layer underneath them. We report this as the better-supported hypothesis, not a fully isolated cause. To check whether the OS-runtime pattern itself is an artifact of one specific GPU, we repeated the identical run on `edu_v100` – an older GPU generation with a slower host interconnect (the whole production run was $\sim 3\times$ slower end-to-end on V100, consistent with H2D transfer time alone rising from $\sim 52\%$ to $\sim 67\%$ of GPU time). Despite the large hardware difference, the OS-runtime split is nearly identical (Figure 12): `fclose` $\sim 17.9\%$, `fwrite` $\sim 0.4\%$, `epoll_wait+poll` $\sim 76\%$ combined on both. This is good evidence the bottleneck is I/O/synchronization behavior, not GPU-specific.

3.7 OpenMP thread scaling and CUDA block size

Figure 13 sweeps OpenMP thread count 1–32 at $M = 4096$ (the full logical-core range of one `edu_a40` rank, $96/3$). Both the full-pipeline and `compute_next`-only curves are essentially flat, never exceeding $\sim 1.2\times$ speedup, with compute-only plateauing by 2 threads and staying flat through 32. This is consistent with the 5-point stencil being memory-bandwidth-bound rather than compute-bound (4




Figure 12: OS Runtime Summary breakdown (rank 0), `edu_a40` vs. `edu_v100` – nearly identical despite very different GPU hardware.

loads, 1 store, few FLOPs per point) – a small number of threads already saturates available memory bandwidth. We also considered an alternative explanation – that `compute_next` forking a fresh OpenMP thread team on every one of 300 calls/rank could plateau this curve through fork/join overhead alone, independent of memory bandwidth – and ruled it out with VTune Hotspots at $M = 4096$: `omp_team_barrier_wait_end` does not appear among the top hotspots at either 4 or 32 threads (unlike part 1’s profile, where it was 79.4% of CPU time), and `compute_next`’s own CPU time is flat across that range (3.43s at 4 threads vs. 3.21s at 32) rather than growing with team size. This contrasts with part 1’s own OpenMP results, which showed real scaling at large thread counts, mainly because part 1 tested M up to 15000 – large enough for the stencil’s absolute cost to be interesting on its own, with no competing GPU-bound stage.




Figure 13: Speedup and efficiency vs. OpenMP thread count at $M = 4096$, full pipeline vs. compute-only.

Figure 14 sweeps the CUDA kernel’s threads-per-block at $M = 4096$: time decreases roughly monotonically as block size grows, worst at 32 threads/block (~ 83 s average) and best at 1024 (~ 73 s, $\sim 12\%$ faster), with only a minor plateau between 128 and 512. Since Nsight Compute (the tool that would report actual occupancy) is unavailable on this cluster, we do not have a profiler-level explanation for this trend, only the timing result itself – the most we can say is that fewer, larger

blocks amortize per-block launch and scheduling overhead across more threads. A smaller, noisier earlier sweep at $M = 1024$ (not shown) trended the opposite way at the high end, so we do not treat this shape as independently confirmed and report only the $M = 4096$ result.




Figure 14: GPU colorize+write time vs. CUDA threads-per-block, $M = 4096$.

3.8 GPU contention across execution modes

Section 3.5 showed disk writes dominate; this raises the question of how much the design choice of sharing one GPU across 3 ranks actually costs. We compared four execution modes at identical parameters ($M = 2048$, 50 steps, 4 threads), isolating sim1’s workload in each: **solo** (1 rank, exclusive GPU), **shared** (the normal 3-rank/1-GPU run), **distributed** (3 ranks, 3 dedicated GPUs on one node, `--gres=gpu:3`), and **multinode** (3 ranks on 3 separate nodes, one dedicated GPU each, communicating over the cluster’s Infiniband fabric [11]). Results (Figure 15): **solo** 9.47s, **shared** 20.84s, **distributed** 21.25s, **multinode** 18.76s.

Two findings. First, dedicating a GPU per rank alone (**distributed** vs. **shared**) makes no reliable difference at all – **distributed** is actually $\sim 2\%$ slower here, well within run-to-run noise rather than a real effect either way – GPU compute-engine time-slicing is not the bottleneck. Second, **solo** is dramatically faster than every multi-rank mode (roughly half the time, $\sim 50\text{--}55\%$ less depending on which mode it’s compared against), which Section 3.5’s finding explains directly: **shared** and **distributed** both have 3 ranks writing frames *concurrently* to the same node; **solo** is the only mode with exactly one writer at a time, anywhere. **multinode** sits in between ($\sim 10\%$ faster than **distributed**) by removing host-local contention (memory bus, PCIe root complex) but not the deeper one: all three nodes still write to the same shared network filesystem (BeeGFS/NFS, cluster-wide by design), so **multinode** is closer to **shared** than to **solo**. The dominant contention in this workload is concurrent I/O, not GPU compute-engine sharing – consistent with Section 3.6’s **fclose**-dominated OS Runtime profile.

3.9 Conclusions

The dominant cost in this pipeline is not GPU compute (colorization kernel: $\sim 1\%$ of frame time) nor CPU compute (**compute_next**: $\sim 2\%$), but per-frame file I/O, specifically whatever happens at **fclose** time rather than the write itself – confirmed consistent across two GPU generations and reflected directly in the GPU-sharing contention results: sharing the GPU across ranks costs almost nothing, while sharing the filesystem costs a great deal. This is the throughline of the whole investigation – every optimization decision was made only after measuring, not assumed from the physics alone. Parallelizing the ASCII-formatting step was considered and rejected once the phase breakdown showed it was $\sim 4\%$ of **color_time**, not worth the added complexity. Per-frame adaptive color scaling was




Figure 15: sim1’s execution time across four GPU/node sharing configurations, identical workload.

adopted only after raw pixel data confirmed a fixed scale reads as pale almost immediately, since a 2D wave’s peak amplitude falls off geometrically with distance from the source on top of γ -damping. And the GPU-sharing question itself was settled by directly comparing four execution modes rather than assumed from GPU specifications – the result (I/O-bound, not compute-bound) held consistently across every measurement in this report.

References

- [1] Politecnico di Torino. Hpc@polito infrastructure guide. https://hpc.polito.it/docs/hpcpolito_guide.pdf.
- [2] Ulrich Drepper. What every programmer should know about memory. Technical report, Red Hat, Inc., 2007.
- [3] Intel Corporation. Hotspots analysis for CPU usage issues – intel vtune profiler user guide. <https://www.intel.com/content/www/us/en/docs/vtune-profiler/user-guide/2026-1/basic-hotspots-analysis.html>.
- [4] Intel Corporation. Threading analysis – intel vtune profiler user guide. <https://www.intel.com/content/www/us/en/docs/vtune-profiler/user-guide/2026-1/threading-analysis.html>.
- [5] Deborah T. Marr, Frank Binns, David L. Hill, Glenn Hinton, David A. Koufaty, J. Alan Miller, and Michael Upton. Hyper-threading technology architecture and microarchitecture. *Intel Technology Journal*, 2002. <https://www.intel.com/content/dam/www/public/us/en/documents/research/2002-vol06-iss-1-intel-technology-journal.pdf>.
- [6] OpenMP Architecture Review Board. OpenMP application programming interface, version 5.0. Technical report, OpenMP Architecture Review Board, November 2018. <https://www.openmp.org/spec-html/5.0/openmp.html>.
- [7] Edgar Ernesto Sanchez Sanchez. 07_cache_mem.pdf course slides, 2026.
- [8] Edgar Ernesto Sanchez Sanchez. 10_parallel_architectures.pdf course slides. 2026.
- [9] Edgar Ernesto Sanchez Sanchez. 14_openmp.pdf course slides. 2026.
- [10] Edgar Ernesto Sanchez Sanchez. 16_mpi_intro.pdf course slides. 2026.
- [11] Edgar Ernesto Sanchez Sanchez. 17_mpi_virtual_topologies.pdf course slides. 2026.
- [12] Edgar Ernesto Sanchez Sanchez. 18_cuda.pdf course slides. 2026.
- [13] Edgar Ernesto Sanchez Sanchez. 19_gpu_programming_intro.pdf course slides. 2026.
- [14] Wikipedia. Finite difference method. https://en.wikipedia.org/wiki/Finite_difference.